



Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

Database Management Systems

Module 43: Indexing and Hashing/3: Indexing/3

Partha Pratim Das

Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

ppd@cse.iitkgp.ac.in



Module 43

Partha Pratim
Das

Objectives & Outline

B+-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary

- Recapitulated the notions of Balanced Binary Search Trees as options for optimal in-memory search data structures
- Understood the issues relating to external data structures for persistent data
- Explored 2-3-4 Tree in depth as a precursor to B/B+-Tree for an efficient external data structure for database and index tables



Module 43

Partha Pratim
Das

Objectives & Outline

B+-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary

- To understand the design of B⁺ Tree Index Files as a generalization of 2-3-4 Tree
- To understand the fundamentals of B-Tree Index Files



Module 43

Partha Pratim
Das

Objectives & Outline

B+-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary

- B⁺ Tree Index Files
- B-Tree Index Files



Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

B⁺ Tree Index Files



Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index


Files

Comparison

Module Summary

The B⁺ Tree

- Is a *balanced binary search tree*
 - Follows a *multi-level index* format like 2-3-4 Tree
- Has the *leaf nodes denoting actual data pointers*
- Ensures that all *leaf nodes remain at the same height* (like 2-3-4 Tree)
- Has the *leaf nodes are linked using a link list*
 - Can support *random access as well as sequential access*
- Example:



Source: *B⁺ Tree*



B⁺ Tree (2)

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys


Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary



- Internal node contains
 - At least $\frac{n}{2}$ child pointers, except the root node
 - At most n pointers
- Leaf node contains
 - At least $\frac{n}{2}$ record pointers and $\frac{n}{2}$ key values
 - At most n record pointer and n key values
 - One block pointer P to point to next leaf node

Source: B⁺ Tree

B⁺ Tree (3): Search

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices


Strings

B-Tree Index
Files

Comparison

Module Summary

- Suppose we have to search 55 in the B⁺ tree below
 - First, we will fetch for the intermediary node which will direct to the leaf node that can contain a record for 55
- So, in the intermediary node, we will find a branch between 50 and 75 nodes
 - Then at the end, we will be redirected to the third leaf node
 - Here DBMS will perform a sequential search to find 55



Source: B⁺ Tree

B⁺ Tree (3): Insert

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices


Strings

B-Tree Index

Files

Comparison

Module Summary



- Suppose we want to insert a record 60 that goes to 3rd leaf node after 55
- The leaf node of this tree is already full, so we cannot insert 60 there
- So we have to split the leaf node, so that it can be inserted into tree without affecting the fill factor, balance and order
- The 3rd leaf node has the values (50, 55, 60, 65, 70) and its current root node is 50
- We will split the leaf node of the tree in the middle so that its balance is not altered

Source: B⁺ Tree

B⁺ Tree (4): Insert

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices


Strings

B-Tree Index

Files

Comparison

Module Summary



- So we can group (50, 55) and (60, 65, 70) into 2 leaf nodes
- If these two has to be leaf nodes, the intermediate node cannot branch from 50
- It should have 60 added to it, and then we can have pointers to a new leaf node
- This is how we can insert an entry when there is overflow. In a normal scenario, it is very easy to find the node where it fits and then place it in that leaf node

Source: [B⁺ Tree](#)

B⁺ Tree (5): Delete

Module 43

Partha Pratim Das

Objectives & Outline

B+-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices


Strings

B-Tree Index Files


Comparison

Comparison

Module Summary



- To delete 60, we have to remove 60 from intermediate node as well as 4th leaf node
- If we remove it from the intermediate node, then the tree will not remain a B+ tree
- So with deleting 60 we re-arranging the nodes:





B⁺ Tree Index Files

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- B⁺ tree indices are an alternative to indexed-sequential files
- *Disadvantage of ISAM files*
 - Performance degrades as file grows, since many overflow blocks get created
 - Periodic reorganization of entire file is required
- *Advantage of B⁺ tree index files:*
 - Automatically reorganizes itself with small, local, changes, in the face of insertions and deletions
 - Reorganization of entire file is not required to maintain performance
- *(Minor) disadvantage of B⁺ trees:*
 - Extra insertion and deletion overhead, space overhead
- *Advantages of B⁺ trees outweigh disadvantages*
 - B⁺ trees are used extensively

B⁺ Tree Index Files (2): Example

Module 43

Partha Pratim Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices

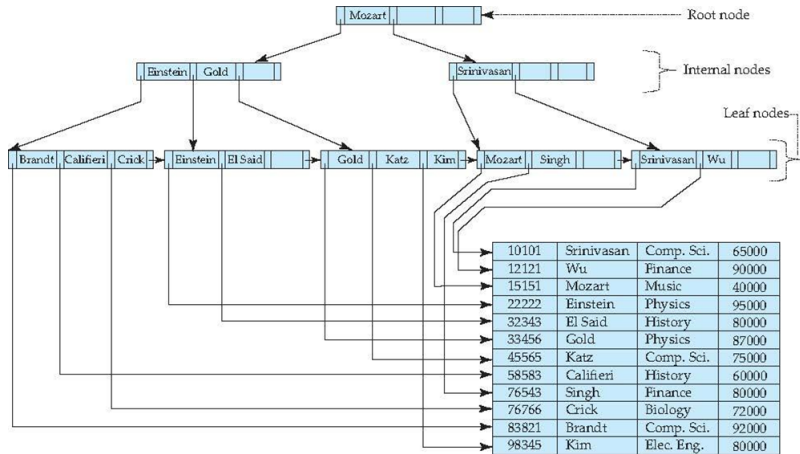
Strings

B-Tree Index

Files

Comparison

Module Summary





B⁺ Tree Index Files (3): Structure

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

A B⁺ tree is a rooted tree satisfying the following properties:

- All paths from root to leaf are of the same length
- Each node that is not a root or a leaf has between $\lceil \frac{n}{2} \rceil$ and n children
- A leaf node has between an $\lceil \frac{n-1}{2} \rceil$ and $n - 1$ values
- Special cases:
 - If the root is not a leaf, it has at least 2 children.
 - If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and $(n - 1)$ values.



B⁺ Tree Index Files (4): Node Structure

Module 43

Partha Pratim Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices


Strings

B-Tree Index Files

Comparison

Module Summary

- Typical node



- K_i are the search-key values
- P_i are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes).
- The search-keys in a node are ordered
 $K_1 < K_2 < K_3 < \dots < K_{n-1}$
(Initially assume no duplicate keys, address duplicates later)



B⁺ Tree Index Files (5): Leaf Nodes

Module 43

Partha Pratim
DasObjectives &
OutlineB+-Tree Index
FilesSimple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

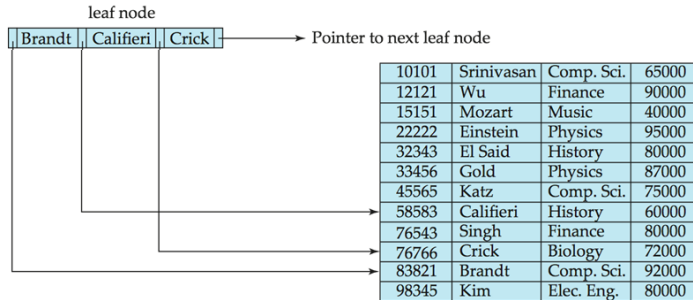
B-Tree Index
Files

Comparison

Module Summary

Properties of a leaf node

- For $i = 1, 2, \dots, n - 1$, pointer P_i points to a file record with search-key value K_i ,
- If L_i, L_j are leaf nodes and $i < j$, L_i 's search-key values are less than or equal to L_j 's search-key values
- P_n points to next leaf node in search-key order





B⁺ Tree Index Files (6): Non-Leaf Nodes

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Non leaf nodes form a multi-level sparse index on the leaf nodes. For a non-leaf node with m pointers:
 - All the search-keys in the subtree to which P_1 points are less than K_1
 - For $2 \leq i \leq n-1$, all the search-keys in the subtree to which P_i points have values greater than or equal to K_{i-1} and less than K_i
 - All the search-keys in the subtree to which P_n points have values greater than or equal to K_{n-1}

P_1	K_1	P_2	...	P_{n-1}	K_{n-1}	P_n
-------	-------	-------	-----	-----------	-----------	-------



B⁺ Tree Index Files (7): Example

Module 43

Partha Pratim
Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices


Strings

B-Tree Index

Files

Comparison

Module Summary



B⁺ tree for *instructor* file ($n = 6$)

- Leaf nodes must have between 3 and 5 values: $\lceil \frac{n-1}{2} \rceil$ and $n-1$, with $n = 6$
- Non-leaf nodes other than root must have between 3 and 6 children: $\lceil \frac{n}{2} \rceil$ and n with $n = 6$
- Root must have at least 2 children



B⁺ Tree Index Files: Observations

Module 43

Partha Pratim
Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary

- Since the inter-node connections are done by pointers, *logically* close blocks need not be *physically* close
- The non-leaf levels of the B⁺ tree form a hierarchy of sparse indices
- The B⁺ tree contains a relatively small number of levels
 - Level below root has at least $2 * \lceil \frac{n}{2} \rceil$ values
 - Next level has at least $2 * \lceil \frac{n}{2} \rceil * \lceil \frac{n}{2} \rceil$ values
 - ... etc.
 - If there are K search-key values in the file, the tree height is no more than $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$
 - thus searches can be conducted efficiently
- Insertions and deletions to the main file can be handled efficiently, as the index can be restructured in logarithmic time

B⁺ Tree Index Files: Queries

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings


B-Tree Index

Files

Comparison

Module Summary

- Find record with search-key value V
 - a) $C = \text{root}$
 - b) While C is not a leaf node
 - i) Let i be least value such that $V \leq K_i$
 - ii) If no such exists, set $C = \text{last non-null pointer in } C$
 - iii) Else $\{ \text{if } (V = K_i) \text{ Set } C = P_{i+1} \text{ else set } C = P_i \}$
 - c) Let i be least value s.t. $K_i = V$
 - d) If there is such a value i , follow pointer P_i to the desired record
 - e) Else no record with search-key value k exists





B⁺ Trees Index Files: Queries (2)

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- If there are K search-key values in the file, the height of the tree is no more than $\left\lceil \log_{\lceil \frac{n}{2} \rceil} (K) \right\rceil$
- A node is generally the same size as a disk block, typically 4 kilobytes
 - and n is typically around 100 (40 bytes per index entry)
- With 1 million search key values and $n = 100$
 - at most $\log_{50}(1,000,000) = 4$ nodes are accessed in a lookup
- Contrast this with a balanced binary tree with 1 million search key values — around 20 nodes are accessed in a lookup
 - above difference is significant since every node access may need a disk I/O, costing around 20 milliseconds



B⁺ Tree Index Files: Handling Duplicates

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- With duplicate search keys
 - In both leaf and internal nodes,
 - ▷ we cannot guarantee that $K_1 < K_2 < K_3 < \dots < K_{n-1}$
 - ▷ but can guarantee $K_1 \leq K_2 \leq K_3 \leq \dots \leq K_{n-1}$
 - Search-keys in the subtree to which P_i points
 - ▷ are $\leq K_i$, but not necessarily $< K_i$,
 - ▷ To see why, suppose same search key value V is present in two leaf node L_i and L_{i+1} . Then in parent node K_i must be equal to V



B⁺ Tree Index Files: Handling Duplicates (2)

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and

Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- We modify find procedure as follows
 - traverse P_i even if $V = K_i$
 - As soon as we reach a leaf node C check if C has only search key values less than V
 - ▷ if so set $C =$ right sibling of C before checking whether C contains V
- Procedure printAll
 - uses modified find procedure to find first occurrence of V
 - Traverse through consecutive leaves to find all occurrences of V



Updates on B⁺ Trees: Insertion

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Find the leaf node in which the search-key value would appear
- If the search-key value is already present in the leaf node
 - Add record to the file
 - If necessary add a pointer to the bucket
- If the search-key value is not present, then
 - Add the record to the main file (and create a bucket if necessary)
 - If there is room in the leaf node, insert (key-value, pointer) pair in the leaf node
 - Otherwise, split the node (along with the new (key-value, pointer) entry) as discussed in the next slide



Updates on B⁺ Trees: Insertion (2)

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings


B-Tree Index

Files

Comparison

Module Summary

- Splitting a leaf node:
 - take the n (search-key value, pointer) pairs (including the one being inserted) in sorted order. Place the first $\lceil \frac{n}{2} \rceil$ in the original node, and the rest in a new node
 - let the new node be p , and let k be the least key value in p . Insert (k, p) in the parent of the node being split
 - If the parent is full, split it and **propagate** the split further up
- Splitting of nodes proceeds upwards till a node that is not full is found
 - In the worst case the root node may be split increasing the height of the tree by 1



Result of splitting node containing Brandt, Califieri and Crick on inserting Adams
Next step: insert entry with (Califieri, pointer-to-new-node) into parent

Updates on B⁺ Trees: Insertion (3)

Module 43

Partha Pratim Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices

Strings



B-Tree Index

Files

Comparison

Module Summary

- Splitting a non-leaf node: when inserting (k, p) into an already full internal node N
 - Copy N to an in-memory area M with space for $n + 1$ pointers and n keys
 - Insert (k, p) into M
 - Copy $P_1, K_1, \dots, K_{\lceil \frac{n}{2} \rceil - 1}, P_{\lceil \frac{n}{2} \rceil}$ from M back into node N
 - Copy $P_{\lceil \frac{n}{2} \rceil + 1}, K_{\lceil \frac{n}{2} \rceil + 1}, \dots, K_n, P_{n+1}$ from M into newly allocated node N'
 - Insert $(K_{\lceil \frac{n}{2} \rceil}, N')$ into parent N
- **Read pseudocode in book!**





Updates on B⁺ Trees: Insertion Example

Module 43

Partha Pratim
Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices



Strings

B-Tree Index

Files

Comparison

Module Summary



B⁺ Tree before and after insertion of "Adams"

Updates on B⁺ Trees: Insertion Example (2)

Module 43

Partha Pratim Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and Secondary Indices

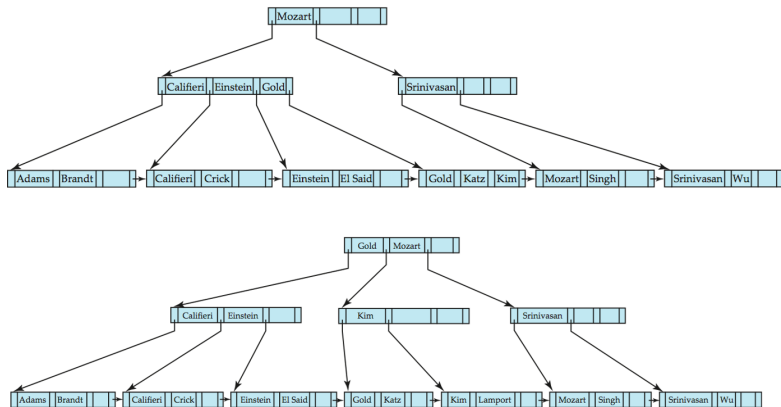
Strings

B-Tree Index

Files

Comparison

Module Summary



B⁺ Tree before and after insertion of "Lampport"



Updates on B⁺ Trees: Deletion

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Find the record to be deleted, and remove it from the main file and from the bucket (if present)
- Remove (search-key value, pointer) from the leaf node if there is no bucket or if the bucket has become empty
- If the node has too few entries due to the removal, and the entries in the node and a sibling fit into a single node, then *merge siblings*:
 - Insert all the search-key values in the two nodes into a single node (the one on the left), and delete the other node.
 - Delete the pair (K_{i-1}, P_i) , where P_i is the pointer to the deleted node, from its parent, recursively using the above procedure.



Updates on B⁺ Trees: Deletion (2)

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Otherwise, if the node has too few entries due to the removal, but the entries in the node and a sibling do not fit into a single node, then **redistribute pointers**:
 - Redistribute the pointers between the node and a sibling such that both have more than the minimum number of entries
 - Update the corresponding search-key value in the parent of the node
- The node deletions may cascade upwards till a node which has $\lceil \frac{n}{2} \rceil$ or more pointers is found
- If the root node has only one pointer after deletion, it is deleted and the sole child becomes the root



Updates on B⁺ Trees: Deletion Example

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices


Strings

B-Tree Index


Files

Comparison

Module Summary



Before and after deleting "Srinivasan"



Deleting "Srinivasan" causes merging of under-full leaves



Updates on B⁺ Trees: Deletion Example (2)

Module 43

Partha Pratim
Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys


Relocation and
Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary



- Leaf containing Singh and Wu became underfull, and borrowed a value Kim from its left sibling
- Search-key value in the parent changes as a result



Updates on B⁺ Trees: Deletion Example (3)

Module 43

Partha Pratim
DasObjectives &
OutlineB+-Tree Index
FilesSimple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

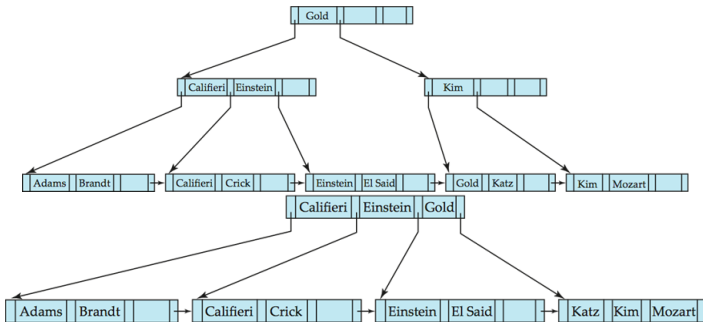
Strings

B-Tree Index

Files

Comparison

Module Summary



Before and after deletion of "Gold" from earlier example

- Node with "Gold" and "Katz" became underfull, and was merged with its sibling
- Parent node becomes underfull, and is merged with its sibling
 - Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is delete

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Index file degradation problem is solved by using B⁺ Tree indices
- Data file degradation problem is solved by using B⁺ Tree File Organization
- The leaf nodes in a B⁺ tree file organization store records, instead of pointers
- Leaf nodes are still required to be half full
 - Since records are larger than pointers, the maximum number of records that can be stored in a leaf node is less than the number of pointers in a non-leaf node
- Insertion and deletion are handled in the same way as insertion and deletion of entries in a B⁺ tree index



B⁺ Tree File Organization: Example

Module 43

Partha Pratim
Das

Objectives & Outline

B⁺-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys


Relocation and Secondary Indices

Strings

B-Tree Index Files

Comparison

Module Summary



Example of B⁺ tree File Organization

- Good space utilization important since records use more space than pointers.
- To improve space utilization, involve more sibling nodes in redistribution during splits and merges
 - Involving 2 siblings in redistribution (to avoid split / merge where possible) results in each node having at least $\lceil \frac{2n}{3} \rceil$ entries



Non-Unique Search Keys

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Alternatives to scheme described earlier
 - Buckets on separate block (bad idea)
 - List of tuple pointers with each key
 - ▷ Extra code to handle long lists
 - ▷ Deletion of a tuple can be expensive if there are many duplicates on search key (why?)
 - ▷ Low space overhead, no extra cost for queries
 - Make search key unique by adding a record-identifier
 - ▷ Extra storage overhead for keys
 - ▷ Simpler code for insertion/deletion
 - ▷ Widely used



Record Relocation and Secondary Indices

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- If a record moves, all secondary indices that store record pointers have to be updated
- Node splits in B⁺ tree file organizations become very expensive
- *Solution:* Use primary-index search key instead of record pointer in secondary index
 - Extra traversal of primary index to locate record
 - Higher cost for queries, but node splits are cheap
 - Add record-id if primary-index search key is non-unique



Indexing Strings

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Variable length strings as keys
 - Variable fanout
 - Use space utilization as criterion for splitting, not number of pointers
- **Prefix compression**
 - Key values at internal nodes can be prefixes of full key
 - ▷ Keep enough characters to distinguish entries in the subtrees separated by the key value
 - For example, “Silas” and “Silberschatz” can be separated by “Silb”
 - Keys in leaf node can be compressed by sharing common prefixes



Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

B-Tree Index Files



B-Tree Index Files

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files


Comparison

Module Summary

- Similar to B⁺ tree, but B-tree allows search-key values to appear only once; eliminates redundant storage of search keys
- Search keys in non-leaf nodes appear nowhere else in the B-tree; an additional pointer field for each search key in a non-leaf node must be included
- Generalized B-tree leaf node



(a)



(b)

- Non-leaf node - pointers B_i are the bucket or file record pointers

B-Tree Index File (2): Example

Module 43

Partha Pratim Das

Objectives & Outline

B+-Tree Index Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys


Relocation and Secondary Indices

Strings


B-Tree Index Files

Comparison

Module Summary



B-tree (above) and B⁺ tree (below) on same data





Comparison of B-Tree and B⁺ Tree Index Files

Module 43

Partha Pratim
Das

Objectives &
Outline

B⁺-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Advantages of B-Tree indices:
 - May use less tree nodes than a corresponding B⁺ Tree
 - Sometimes possible to find search-key value before reaching leaf node
- Disadvantages of B-Tree indices:
 - Only small fraction of all search-key values are found early
 - Non-leaf nodes are larger, so fan-out is reduced. Thus, B-Trees typically have greater depth than corresponding B⁺ Tree
 - Insertion and deletion more complicated than in B⁺ Trees
 - Implementation is harder than B⁺ Trees
- Typically, advantages of B-Trees do not outweigh disadvantages



Module Summary

Module 43

Partha Pratim
Das

Objectives &
Outline

B+-Tree Index
Files

Simple B⁺ Tree

Index Files

Nodes

Observations

Query

Duplicates

Updates

Insertion

Deletion

File Organization

Non-Unique Keys

Relocation and
Secondary Indices

Strings

B-Tree Index
Files

Comparison

Module Summary

- Understood the design of B⁺ Tree Index Files in depth for database persistent store
- Familiarized with B-Tree Index Files

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.