



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Database Management Systems

Module 56: Query Processing and Optimization/1: Processing

Partha Pratim Das

Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

ppd@cse.iitkgp.ac.in



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Learnt the importance of backup and analysed different backup strategies
- Failures may be due to variety of sources – each needs a strategy for handling
- A proper mix and management of volatile, non-volatile and stable storage can guarantee recovery from failures and ensure Atomicity, Consistency and Durability
- Log-based recovery is efficient and effective
- Learnt how Hot backup of transaction log helps in recovering consistent database.
- Studied the recovery algorithms for concurrent transactions
- Recovery based on operation logging supplements log-based recovery
- Planning for Backup
- Understood RAID - array of redundant disks in parallel to enhance speed and reliability



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- To understand the overall flow for Query Processing
- To define the Measures of Query Cost
- To understand the algorithms for processing Selection Operations, Sorting, Join Operations, and a few Other Operations



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Overview of Query Processing
- Measures of Query Cost
- Selection Operation
- Sorting
- Join Operation
- Other Operations



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

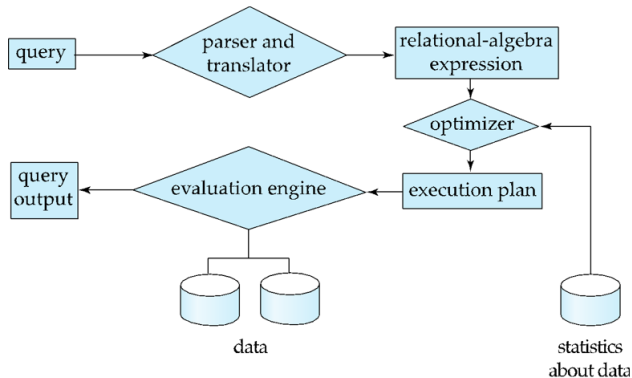
Module Summary

Overview of Query Processing



Basic Steps in Query Processing

- a) Parsing and translation
- b) Optimization
- c) Evaluation





Basic Steps in Query Processing (2)

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Parsing and translation
 - translate the query into its internal form
 - ▷ This is then translated into relational algebra
 - Parser checks syntax, verifies relations
- Evaluation
 - The query-execution engine takes a query-evaluation plan, executes that plan, and returns the answers to the query



Basic Steps in Query Processing (3): Optimization

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Consider the query

```
select salary  
from instructor  
where salary < 75000;
```

which can be translated into either of the following relational-algebra expressions:

- $\sigma_{salary < 75000}(\Pi_{salary}(instructor))$
- $\Pi_{salary}(\sigma_{salary < 75000}(instructor))$
- Each relational algebra operation can be evaluated using one of several different algorithms
 - Correspondingly, a relational-algebra expression can be evaluated in many ways
- Annotated expression specifying detailed evaluation strategy is called an **evaluation-plan**.
 - For example, can use an index on salary to find instructors with salary < 75000,
 - or can perform complete relation scan and discard instructors with salary ≥ 75000



Basic Steps in Query Processing (4): Optimization

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Query Optimization:** Amongst all equivalent evaluation plans choose the one with lowest cost
 - Cost is estimated using statistical information from the database catalog
 - ▷ For example, number of tuples in each relation, size of tuples, etc.
- In this module we study
 - How to measure query costs
 - Algorithms for evaluating relational algebra operations
 - How to combine algorithms for individual operations in order to evaluate a complete expression
- In the next module
 - We study how to optimize queries, that is, how to find an evaluation plan with lowest estimated cost

Measures of Query Cost

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Measures of Query Cost



Measures of Query Cost

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Cost is generally measured as total elapsed time for answering query
 - Many factors contribute to time cost
 - ▷ *disk accesses, CPU, or even network communication*
- Typically disk access is the predominant cost, and is also relatively easy to estimate
- Measured by taking into account
 - Number of seeks * average-seek-cost
 - Number of blocks read * average-block-read-cost
 - Number of blocks written * average-block-write-cost
 - ▷ Cost to write a block is greater than cost to read a block
 - data is read back after being written to ensure that the write was successful



Measures of Query Cost (2)

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- For simplicity we just use the **number of block transfers** *from disk and the number of seeks* as the cost measures
 - t_T : time to transfer one block
 - t_S : time for one seek
 - Cost for b block transfers plus S seeks

$$b * t_T + S * t_S$$

- We ignore CPU costs for simplicity
 - Real systems do take CPU cost into account
- We do not include cost to writing output to disk in our cost formulae



Measures of Query Cost (3)

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Several algorithms can reduce disk IO by using extra buffer space
 - Amount of real memory available to buffer depends on other concurrent queries and OS processes, known only during execution
 - ▷ We often use worst case estimates, assuming only the minimum amount of memory needed for the operation is available
- Required data may be buffer resident already, avoiding disk I/O
 - But hard to take into account for cost estimation



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

**Selection
Operation**

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Selection Operation



Selection Operation: File / Index Scan

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

A#	Algorithm	Cost	Reason
A1	Linear Search	$t_S + b_r \times t_T$	One initial seek plus b_r block transfers
A1	Linear Search, Eq. on Key	Average case $t_S + (b_r/2) \times t_T$	Since at most one record satisfies condition, scan can be terminated as soon as the required record is found. b_r blocks transfers in worst case
A2	Prm. Index, Eq. on Key	$(h_i + 1) \times (t_T + t_S)$	Index lookup traverses the height of the tree plus one I/O to fetch the record; each of these I/O operations requires a seek and a block transfer
A3	Prm. Index, Eq. on Nonkey	$h_i \times (t_T + t_S) + b \times t_T$	One seek for each level of the tree, one seek for the first block. Here all of b are read. These blocks are leaf blocks assumed to be stored sequentially (for a primary index) and don't require additional seeks
A4	Snd. Index, Eq. on Key	$(h_i + 1) \times (t_T + t_S)$	This case is similar to primary index
A4	Snd. Index, Eq. on Nonkey	$(h_i + n) \times (t_T + t_S)$	Here, cost of index traversal is the same as for A3, but each record may be on a different block, requiring a seek per record. Cost is potentially very high if n is large
A5	Prm. Index, Comparison	$h_i \times (t_T + t_S) + b \times t_T$	Identical to the case of A3, equality on nonkey
A6	Snd. Index, Comparison	$(h_i + n) \times (t_T + t_S)$	Identical to the case of A4, equality on nonkey

t_T is time to transfer one block. t_S is time for one seek

b_r denotes the number of blocks in the file

b denotes the number of blocks containing records with the specified search key

h_i denotes the height of the index. n is the number of records fetched



Complex Selections: Conjunction

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Conjunction:** $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \theta_n}(r)$
- **A7 (conjunctive selection using one index)**
 - Select a combination of θ_i and algorithms A1 through A6 that results in the least cost for $\sigma_{\theta_i}(r)$
 - Test other conditions on tuple after fetching it into memory buffer
- **A8 (conjunctive selection using composite index)**
 - Use appropriate composite (multiple-key) index if available
- **A9 (conjunctive selection by intersection of identifiers)**
 - Requires indices with record pointers
 - Use corresponding index for each condition, and take intersection of all the obtained sets of record pointers
 - Then fetch records from file
 - If some conditions do not have appropriate indices, apply test in memory



Complex Selections: Disjunction

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Disjunction:** $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$.
- **A10 (disjunctive selection by union of identifiers)**
 - Applicable if *all* conditions have available indices
 - ▷ Otherwise use linear scan
 - Use corresponding index for each condition, and take union of all the obtained sets of record pointers
 - Then fetch records from file
- **Negation:** $\sigma_{\neg\theta}(r)$
 - Use linear scan on file
 - If very few records satisfy $\neg\theta$, and an index is applicable to θ
 - ▷ Find satisfying records using index and fetch from file



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Sorting



Sorting

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- We may **build an index on the relation**, and then use the index to read the relation in sorted order
 - May lead to one disk block access for each tuple
- For relations that fit in memory, techniques like **quicksort** can be used
- For relations that do not fit in memory, **external sort-merge** is a good choice



External Sort-Merge: Example

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

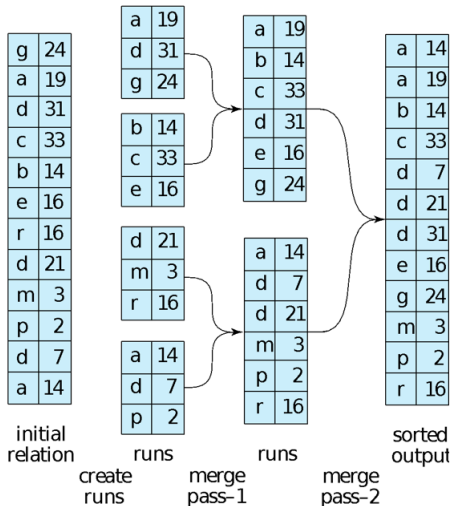
Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary





External Sort-Merge: Algorithm

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- a) **Create sorted runs.** Let M denote the number of blocks in the main-memory buffer available for sorting. First, a number of sorted runs are created; each run is sorted, but contains only some of the records of the relation.

```

i = 0;
repeat
    read M blocks of the relation, or the rest of the relation, whichever is smaller;
    sort the in-memory part of the relation;
    write the sorted data to run file  $R_i$ ;
    i = i + 1;
until the end of the relation

```

- b) **Merge the runs (N-way merge):** Now, the runs are merged. For the total number of runs, $N < M$, so that we can allocate one block to each run and have space left to hold one block of output. The merge stage operates as follows:

```

read one block of each of the  $N$  files  $R_i$  into a buffer block in memory;
repeat
    choose the first tuple (in sort order) among all buffer blocks;
    write the tuple to the output, and delete it from the buffer block;
    if the buffer block of any run  $R_i$  is empty and not end-of-file( $R_i$ )
        then read the next block of  $R_i$  into the buffer block;
until all input buffer blocks are empty

```

- c) If $N \geq M$, several merge passes are required
- In each pass, contiguous groups of $M-1$ runs are merged.
 - A pass reduces the number of runs by a factor of $M-1$, and creates runs longer by the same factor
 - For $M=11$ and 90 runs, one pass reduces the number of runs to 9, each 10 times the size of the initial runs
 - Repeated passes are performed till all runs have been merged into one



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Join Operation



Join Operation

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Several different algorithms to implement joins
 - **Nested-loop join**
 - **Block nested-loop join**
 - **Indexed nested-loop join**
 - Merge-join
 - Hash-join
- Choice based on cost estimate
- Examples use the following information
 - Number of records of *student*: $n_{students} = 5,000$
 - Number of records of *takes*: $n_{takes} = 10,000$
 - Number of blocks of *student*: $b_{students} = 100$
 - Number of blocks of *takes*: $b_{takes} = 400$



Nested-Loop Join

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- To compute the theta join $r \bowtie_{\theta} s$
 for each tuple t_r **in** r **do begin**
 for each tuple t_s **in** s **do begin**
 test pair (t_r, t_s) to see if they satisfy the join condition θ
 if they do, add $t_r \bullet t_s$ to the result.
 end
 end
- r is called the **outer relation** and s the **inner relation** of the join
- Requires no indices and can be used with any kind of join condition
- Expensive since it examines every pair of tuples in the two relations



Nested-Loop Join (2)

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- In the worst case, if there is enough memory only to hold one block of each relation, the estimated cost is $n_r * b_s + b_r$ block transfers, plus $n_r + b_r$ seeks, where n_r (n_s) denotes the number of tuples in r (s) and b_r (b_s) denotes the number of blocks containing tuples of in r (s)
- If the smaller relation fits entirely in memory, use that as the inner relation.
 - Reduces cost to $b_r + b_s$ block transfers and 2 seeks
- Example of join of *students* and *takes*: $n_{students} = 5,000$, $n_{takes} = 10,000$, $b_{students} = 100$, $b_{takes} = 400$
- Assuming worst case memory availability cost estimate is
 - with *student* as outer relation:
 - ▷ $5000 * 400 + 100 = 2,000,100$ block transfers,
 - ▷ $5000 + 100 = 5100$ seeks
 - with *takes* as the outer relation
 - ▷ $10000 * 100 + 400 = 1,000,400$ block transfers and 10,400 seeks
- If smaller relation (*student*) fits entirely in memory, the cost estimate will be 500 block transfers
- Block nested-loops algorithm is preferable



Block Nested-Loop Join

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation

```
for each block  $B_r$  of  $r$  do begin  
  for each block  $B_s$  of  $s$  do begin  
    for each tuple  $t_r$  in  $B_r$  do begin  
      for each tuple  $t_s$  in  $B_s$  do begin  
        Check if  $(t_r, t_s)$  satisfy the join condition  
        if they do, add  $t_r \bullet t_s$  to the result.  
      end  
    end  
  end  
end
```



Block Nested-Loop Join (2)

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Worst case estimate: $b_r * b_s + b_r$ block transfers + $2 * b_r$ seeks
 - Each block in the inner relation s is read once for each block in the outer relation
- Best case: $b_r + b_s$ block transfers + 2 seeks.
- Improvements to nested loop and block nested loop algorithms:
 - In block nested-loop, use $M - 2$ disk blocks as blocking unit for outer relations, where M = memory size in blocks; use remaining two blocks to buffer inner relation and output
 - ▷ $Cost = \lceil b_r / (M - 2) \rceil * b_s + b_r$ block transfers + $2 * \lceil b_r / (M - 2) \rceil$ seeks
 - If equi-join attribute forms a key or inner relation, stop inner loop on first match
 - Scan inner loop forward and backward alternately, to make use of the blocks remaining in buffer (with LRU replacement)
 - Use index on inner relation, if available



Indexed Nested-Loop Join

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Index lookups can replace file scans if
 - join is an equi-join or natural join and
 - an index is available on the inner relation's join attribute
 - ▷ Can construct an index just to compute a join.
- For each tuple t_r in the outer relation r , use the index to look up tuples in s that satisfy the join condition with tuple t_r .
- Worst case: buffer has space for only one page of r , and, for each tuple in r , we perform an index lookup on s .
- Cost of the join: $b_r \times (t_T + t_S) + n_r * c$
 - Where c is the cost of traversing index and fetching all matching s tuples for one tuple or r
 - c can be estimated as cost of a single selection on s using the join condition.
- If indices are available on join attributes of both r and s , use the relation with fewer tuples as the outer relation.



Example of Nested-Loop Join Costs

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Compute $student \bowtie takes$, with $student$ as the outer relation.
- Let $takes$ have a primary B^+ -tree index on the attribute ID , which contains 20 entries in each index node.
- Since $takes$ has 10,000 tuples, the height of the tree is 4, and one more access is needed to find the actual data
- $student$ has 5000 tuples
- Cost of block nested loops join
 - $400 \times 100 + 100 = 40,100$ block transfers + $2 \times 100 = 200$ seeks
 - ▷ assuming worst case memory
 - ▷ may be significantly less with more memory
- Cost of indexed nested loops join
 - $100 + 5000 \times 5 = 25,100$ block transfers and seeks.
 - CPU cost likely to be less than that for block nested loops join



Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

Other Operations



Other Operations

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Duplicate Elimination**
- **Projection**
- **Aggregation**
- Set Operations
- Outer Join



Other Operations: Duplicate Elimination & Projection

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Duplicate Elimination** can be implemented via hashing or sorting
 - On sorting duplicates will come adjacent to each other, and all but one set of duplicates can be deleted
 - *Optimization*: duplicates can be deleted during run generation as well as at intermediate merge steps in external sort-merge
 - Hashing is similar – duplicates will come into the same bucket
- **Projection** :
 - perform projection on each tuple
 - followed by duplicate elimination



Other Operations: Aggregation

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- **Aggregation** can be implemented in a manner similar to duplicate elimination
 - Sorting or hashing can be used to bring tuples in the same group together, and then the aggregate functions can be applied on each group
 - *Optimization*: combine tuples in the same group during run generation and intermediate merges, by computing partial aggregate values
 - ▷ For count, min, max, sum: keep aggregate values on tuples found so far in the group
 - When combining partial aggregate for count, add up the aggregates
 - ▷ For avg, keep sum and count, and divide sum by count at the end



Module Summary

Module 56

Partha Pratim
Das

Week Recap

Objectives &
Outline

Query Processing

Query Cost

Selection
Operation

Complex Selections

Sorting

External Sort-Merge

Join Operation

Other Operations

Module Summary

- Understood the overall flow for Query Processing and defined the Measures of Query Cost
- Studied the algorithms for processing Selection Operations, Sorting, Join Operations and a few Other Operations

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.